

NAME

vsearch — dereplicate, filter, sort, search, compare and clusterize amplicons from metagenomic projects

SYNOPSIS

vsearch [*options*] *filename*

DESCRIPTION

Environmental or clinical molecular studies generate large volumes of amplicons (e.g. SSU-rRNA sequences) that need to be filtered, dereplicated, searched, clustered or compared to sequences from other studies. The aim of **vsearch** is to offer a all-in-one open source tool to perform these tasks, using optimized algorithm implementations and harvesting the full potential of modern computers to guarantee the fastest and more accurate possible processing.

Nucleotidic sequence comparisons is at the core of **vsearch**. To speed up comparisons, **vsearch** implements *k*-mer filtering, and an extremely fast Needleman-Wunsch algorithm making use of the Streaming SIMD Extensions (SSE2) of modern x86-64 CPUs. If SSE2 instructions are not available, **vsearch** exits with an error message.

Input

vsearch input is a fasta file containing one or several nucleotidic sequences. For each sequence, the sequence identifier is defined as the string comprised between the ">" symbol and the first space or the end of the line, whichever comes first. Additionally, if the line ends with the pattern ";size=*integer*", **vsearch** will interpret *integer* as the abundance of the sequence (in a dereplicated fasta file for instance). The nucleotidic sequence is defined as a string of [acgt] or [acgu] symbols (case insensitive), starting after the end of the identifier line and ending before the next identifier line or the file end; **vsearch** exits with an error message if any other symbol is present in the sequence. Optionally, **vsearch** can be compiled to accepted compressed fasta files as input (gz and bzip2 formats).

Options

vsearch recognizes a large number of command-line options. For an easier navigation, options are grouped by theme (dereplication, filtering, sorting, searching, comparison, clustering). We start with general options that apply to all themes.

General options:

- help** display a short help and exit.
- version** output version information and exit.
- fasta_width** *positive integer*
fasta files produced by **vsearch** are wrapped (sequences written on lines of *integer* nucleotides, 80 by default). Set that value to 0 to eliminate the wrapping.
- maxseqlength** *positive integer*
all **vsearch** operations will discard sequences of length equal or greater than *integer* (50,000 nucleotides by default).
- minseqlength** *positive integer*
all **vsearch** operations will discard sequences of length smaller than *integer* (1 nucleotide by default for sorting or shuffling, 32 nucleotides for dereplication or searching).
- notrunclabels**
do not truncate sequence labels at first space, use the full header.
- strand** *plus/both*
when searching or dereplicating, check the *plus* strand only (default) or check *both* strands.
- threads** *positive integer*
number of computation threads to use. The number of threads should be lesser or equal to the number of available CPU cores. The default is to launch one thread per available CPU core.

- uc** *filename*
when searching, clustering or dereplicating, output results in *filename* using a uclust-like format.
- uc_allhits**
when searching, clustering or dereplicating, and when using the uclust-like format option, show all hits, not just top hit.

Clustering options:

- centroids** *filename*
output cluster centroid sequences to *filename* file.
- cluster_smallmem** *filename*
use a slower clustering algorithm (consumes less memory) and write the results to *filename*.
- clusters** *string*
output each cluster to a separate fasta file using *string* as prefix for the filenames.
- consout** *filename*
output cluster consensus sequences to *filename* file.
- construncate**
when using the consout option, do not ignore terminal gaps in the multiple sequence alignment when building the consensus sequence.
- msaout** *filename*
output multiple sequence alignments to *filename*.
- uc** *filename*
use the fast clustering algorithm and write the results to *filename*.
- usersort**
conserve the initial input order of sequences, do not sort sequences by decreasing length.

Dereplication, masking, shuffling and sorting options:

- derep_fulllength** *filename*
merge strictly identical sequences contained in *filename*. Redundant sequences receive the header of the sequence of their group, and the number of occurrences (abundance) is indicated at the end of the fasta header using the pattern ";size=X;".
- maskfasta** *filename*
mask sequences contained in *filename*.
- maxsize** *positive integer*
when using --sortbysize, discard sequences with an abundance value greater than *integer*.
- minsize** *positive integer*
when using --sortbysize, discard sequences with an abundance value smaller than *integer*.
- minuniquesize** *positive integer*
when dereplicating, discard sequences with an abundance value smaller than *integer*.
- output** *filename*
when dereplicating, sorting or shuffling, write the results to *filename*.
- relabel** *string*
when sorting, relabel sequence headers using *string* as suffix.

- seed** *positive integer*
when shuffling, use *integer* as seed. Set to 0 to use a pseudo-random seed.
- sizein** read abundance annotation from input
- sizeout** add abundance annotation to output
- shuffle** *filename*
pseudo-randomly shuffle the order of sequences contained in *filename*.
- sortbylength** *filename*
sort by decreasing length the sequences contained in *filename*.
- sortbysize** *filename*
sort by decreasing abundance the sequences contained in *filename*.
- topn** *positive integer*
when dereplicating, sorting or shuffling, output just the top *integer* sequences.

Searching options:

- alnout** *filename*
write pairwise global alignments to *filename* using a human-readable format.
- blast6out** *filename*
write search results to *filename* using a blast-like tab-separated format.
- db** *filename*
compare query sequences to the fasta-formatted subject sequences contained in *filename*, using global pairwise alignment.
- dbmask** *none/dust/soft*
mask simple repeats and low-complexity regions in subject database sequences using the *dust* or the *soft* algorithms, or do not mask (*none*). The default is to mask using *dust*.
- dbmatched** *filename*
write database subject sequences matching at least one query sequence to *filename*, in fasta format.
- dbnotmatched** *filename*
write database subject sequences not matching query sequences to *filename*, in fasta format.
- fastapairs** *filename*
write pairs of query and subject sequences to *filename*, in fasta format.
- fulldp** uses a 8-way SIMD vectorized full dynamic programming algorithm (Needleman-Wunsch). That option increases the sensitivity of vsearch.
- gapext** *string*
penalties for gap extension (2I/1E)
- gapopen** *string*
penalties for gap opening (20I/2E)
- hardmask**
mask low-complexity regions by replacing them with Ns instead of setting them to lower case.
- id** *real* reject the sequence match if the pairwise identity is lower than *real* (value ranging from 0.0 to 1.0 included).
- idprefix** *positive integer*
reject the subject sequence if the first *integer* nucleotides do not match the query sequence.

- idsuffix** *positive integer*
reject the subject sequence if the last *integer* nucleotides do not match the query sequence.
- leftjust** reject the subject sequence if the alignment begins with gaps.
- leftjust** reject the subject sequence if the alignment begins with gaps.
- match** *integer*
score assigned to a match (i.e. identical nucleotides) in the pairwise alignment. The default value is 2.
- matched** *filename*
write query sequences matching database subject sequences to *filename*, in fasta format.
- maxaccepts** *positive integer*
maximum number of hits to accept before stopping the search. The default value is 1. That option works in pair with maxrejects. The search process sorts subject sequences by decreasing number of kmers they have in common with the query sequence, using that information as a proxy for sequence similarity. If the first subject sequence passes the acceptance criteria, it is accepted as best hit and the search process stops for that query. If maxaccepts is set to a higher value, it accepts more hits. If maxaccepts and maxrejects are both set to 0, the complete database is searched.
- maxdiffs** *positive integer*
reject the subject sequence if the alignment contains at least *integer* substitutions, insertions or deletions.
- maxgaps** *positive integer*
reject the subject sequence if the alignment contains at least *integer* insertions or deletions.
- maxhits** *positive integer*
maximum number of hits to show. The default value is 1.
- maxid** *real*
reject the subject sequence if its percentage of identity with the query is equal or greater than *real*.
- maxqsize** *positive integer*
reject query sequences with an abundance equal or greater than *integer*.
- maxqt** *real*
reject if the query/subject length ratio is equal or greater than *real*.
- maxrejects** *positive integer*
maximum number of non-matching subject sequences to consider before stopping the search. The default value is 32. That option works in pair with maxaccepts. The search process sorts subject sequences by decreasing number of kmers they have in common with the query sequence, using that information as a proxy for sequence similarity. If none of the first 32 subject sequence pass the acceptance criteria, the search process stops for that query (no hit). If maxrejects is set to a higher value, more subject sequences will be considered. If maxaccepts and maxrejects are both set to 0, the complete database is searched.
- maxsizratio** *real*
reject if the query/subject abundance ratio is equal or greater than *real*.
- maxsl** *real*
reject if the shorter/longer length ratio is equal or greater than *real*.

- maxsubs** *positive integer*
reject the subject sequence if the alignment contains at least *integer* substitutions.
- mid** *real*
reject the subject sequence if its percentage of identity with the query is lower than *real* (ignoring gaps).
- mincols** *positive integer*
reject the subject sequence if the alignment length is shorter than *integer*.
- minqt** *real*
reject if the query/subject length ratio is lower than *real*.
- minsizeratio** *real*
reject if the query/subject abundance ratio is lower than *real*.
- minsl** *real*
reject if the shorter/longer length ratio is lower than *real*.
- mintsiz** *positive integer*
reject subject sequences with an abundance lower than *integer*.
- mismatch** *integer*
score assigned to a mismatch (i.e. different nucleotides) in the pairwise alignment. The default value is -4.
- notmatched** *filename*
write query sequences not matching database subject sequences to *filename*, in fasta format.
- output_no_hits** *filename*
write both matching and non-matching queries to output files. Non-matching queries are labelled "no hit" (**to be verified**).
- qmask** *none/dust/soft*
mask simple repeats and low-complexity regions in query sequences using the *dust* or the *soft* algorithms, or do not mask (*none*). The default is to mask using *dust*.
- query_cov** *real*
reject if the fraction of the query aligned to the subject sequence is lower than *real*.
- rightjust**
reject the subject sequence if the alignment ends with gaps.
- rowlen** *positive integer*
width of alignment lines in alnout output. The default value is 64.
- self**
reject the alignment if the query and subject labels are identical.
- selfid**
reject the alignment if the query and subject sequences are identical.
- target_cov** *real*
reject if the fraction of the subject sequence aligned to the query sequence is lower than *real*.
- top_hits_only**
output only hits with the highest percentage of identity with the query.
- userfields** *string*
when using the userout option, select the fields that are written to the file. (**list and document the fields**)
- userout** *filename*
write user-defined tab-separated output to *filename*. See "userfields".

--vsearch_global *filename*

filename of queries for global alignment search.

--weak_id *real*

show hits with percentage of identity of at least *real*; and continue the search. That option allows to report very weak sequence similarities.

--wordlength *positive integer*

length of words (*k*mers) for database index. The default value is 8.

EXAMPLES

(in progress)

Search queries in a reference database:

```
vsearch --vsearch_global queries.fas --db references.fas --alnout results.aln --id 0.8 --fuldp
```

search a sequence dataset against itself, ignoring self hits but getting all matches above a certain identity level, and collecting results in a blast-like format:

```
vsearch --vsearch_global queries.fas --db queries.fas --strand plus --id 0.6 --alnout results.aln
--selfid --blast6out results.blast6 -maxaccepts 0 -maxrejects 0
```

LIMITATIONS

vsearch does not yet perform chimera detection.

AUTHORS

Implementation by Torbjørn Rognes and Tomas Flouri, documentation by Frédéric Mahé, .

REPORTING BUGS

Submit suggestions and bug-reports at <<https://github.com/torognes/vsearch/issues>>, send a pull request on <<https://github.com/torognes/vsearch>>, or compose a friendly or curmudgeont e-mail to Torbjørn Rognes <torognes@ifi.uio.no>.

AVAILABILITY

The software is available from <<https://github.com/torognes/vsearch>>

COPYRIGHT

Copyright (C) 2014 Torbjørn Rognes et al.

This program is free software: you can redistribute it and/or modify it under the terms of the GNU Affero General Public License as published by the Free Software Foundation, either version 3 of the License, or any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Affero General Public License for more details.

You should have received a copy of the GNU Affero General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

vsearch includes code from Google's CityHash project by Geoff Pike and Jyrki Alakuijala, providing some excellent hash functions available under a MIT license.

vsearch includes code derived from Tatusov and Lipman's DUST program that is in the public domain.

vsearch binaries may include code from the zlib library copyright Jean-loup Gailly and Mark Adler.

vsearch binaries may include code from the bzip2 library copyright Julian R. Seward.

SEE ALSO

swipe, an extremely fast Smith-Waterman database search tool by Torbjørn Rognes (available from [<https://github.com/torognes/swipe>](https://github.com/torognes/swipe)).

VERSION HISTORY

New features and important modifications of **vsearch** (short lived or minor bug releases are not mentioned):

v1.0 released November 1st, 2014

First public release